Research Analysis Report

Research Topic: Multi-Agent Systems

Total Papers: 5

Generated By: Report Generation Agent

Generated At: 2026-08-27 14:14:50

Executive Summary

This report provides a technical overview of 'Multi-Agent Systems' including methodologies, research gaps and experiment planning.

Literature Review

Introduction: This literature review summarizes findings from 5 research papers related to the selected topic. It highlights research trends, methodologies, key contributions, future directions, and overall observations.

Paper Summaries

• {'title': 'Consensus and Cooperation in Networked Multi-Agent Systems', 'year': 2007, 'summary': 'Abstract not available for analysis.'}

• {'title': 'Why Do Multi-Agent LLM Systems Fail?', 'year': 2025, 'summary': 'Despite enthusiasm for Multi-Agent LLM Systems (MAS), their performance gains on popular benchmarks are often minimal. This gap highlights a critical need for a principled understanding of why MAS fail. Addressing this question requires systematic identification and analysis of failure patterns. We introduce MAST-Data, a comprehensive dataset of 1600+ annotated traces collected across 7 popular MAS frameworks. MAST-Data is the first multi-agent system dataset to outline the failure dynamics in MAS for guiding the development of better future systems. To enable systematic classification of failures for MAST-Data, we build the first Multi-Agent System Failure Taxonomy (MAST). We develop MAST through rigorous analysis of 150 traces, guided closely by expert human annotators and validated by high inter-annotator agreement (kappa = 0.88). This process identifies 14 unique modes, clustered into 3 categories: (i) system design issues, (ii) inter-agent misalignment, and (iii) task verification. To enable scalable annotation, we develop an LLM-as-a-Judge pipeline with high agreement with human annotations. We leverage MAST and MAST-Data to analyze failure patterns across models (GPT4, Claude 3, Qwen2.5, CodeLlama) and tasks (coding, math, general agent), demonstrating improvement headrooms from better MAS design. Our analysis provides insights revealing that identified failures require more sophisticated solutions, highlighting a clear'}

• {'title': 'Which Agent Causes Task Failures and When? On Automated Failure Attribution of LLM Multi-Agent Systems', 'year': 2025, 'summary': "Failure attribution in LLM multi-agent systems-identifying the agent and step responsible for task failures-provides crucial clues for systems debugging but remains underexplored and labor-intensive. In this paper, we propose and formulate a new research area: automated failure attribution for LLM multi-agent systems. To support this initiative, we introduce the Who&When; dataset, comprising extensive failure logs from 127 LLM multi-agent systems with fine-grained annotations linking failures to specific agents and decisive error steps. Using the Who&When;, we develop and evaluate three automated failure attribution methods, summarizing their corresponding pros and cons. The best method achieves

53.5% accuracy in identifying failure-responsible agents but only 14.2% in pinpointing failure steps, with some methods performing below random. Even SOTA reasoning models, such as OpenAI o1 and DeepSeek R1, fail to achieve practical usability. These results highlight the task's complexity and the need for further research in this area. Code and dataset are available at https://github.com/mingyin1/Agents_Failure_Attribution"}

- {'title': 'Multi-Agent Actor-Critic for Mixed Cooperative-Competitive Environments', 'year': 2017, 'summary': 'We explore deep reinforcement learning methods for multi-agent domains. We begin by analyzing the difficulty of traditional algorithms in the multi-agent case: Q-learning is challenged by an inherent non-stationarity of the environment, while policy gradient suffers from a variance that increases as the number of agents grows. We then present an adaptation of actor-critic methods that considers action policies of other agents and is able to successfully learn policies that require complex multi-agent coordination. Additionally, we introduce a training regimen utilizing an ensemble of policies for each agent that leads to more robust multi-agent policies. We show the strength of our approach compared to existing methods in cooperative as well as competitive scenarios, where agent populations are able to discover various physical and informational coordination strategies.'}
- {'title': 'Red-Teaming LLM Multi-Agent Systems via Communication Attacks', 'year': 2025, 'summary': 'Large Language Model-based Multi-Agent Systems (LLM-MAS) have revolutionized complex problem-solving capability by enabling sophisticated agent collaboration through message-based communications. While the communication framework is crucial for agent coordination, it also introduces a critical yet unexplored security vulnerability. In this work, we introduce Agent-in-the-Middle (AiTM), a novel attack that exploits the fundamental communication mechanisms in LLM-MAS by intercepting and manipulating inter-agent messages. Unlike existing attacks that compromise individual agents, AiTM demonstrates how an adversary can compromise entire multi-agent systems by only manipulating the messages passing between agents. To enable the attack under the challenges of limited control and role-restricted communication format, we develop an LLM-powered adversarial agent with a reflection mechanism that generates contextually-aware malicious instructions. Our comprehensive evaluation across various frameworks, communication structures, and real-world applications demonstrates that LLM-MAS is vulnerable to communication-based attacks, highlighting the need for robust security measures in multi-agent systems.'}

Research Trends

- able
- achieve
- actor-critic
- adaptation
- additionally
- agent
- agent-in-the-middle
- agreement
- annotation
- area
- attack
- attribution
- automated
- benchmark
- better

- challenged
- communication
- communication-based
- competitive
- contextually-aware
- cooperative
- coordination
- dataset
- design
- difficulty
- failure
- failure-responsible
- failures-provide
- fine-grained
- inter-agent
- inter-annotator
- labor-intensive
- llm-as-a-judge
- llm-ma
- llm-powered
- mast
- mast-data
- message-based
- model-based
- multi-agent
- non-stationarity
- policy
- problem-solving
- q-learning
- real-world
- reasoning
- reinforcement
- role-restricted
- steps
- systematic
- systems-identifying

Research Gaps

• These results highlight the task's complexity and the need for further research in this area.

Future Scope

• These results highlight the task's complexity and the need for further research in this area.

Conclusion: Overall, the reviewed studies demonstrate significant progress in this research area while highlighting several opportunities for future work and further experimentation.

Methodology Comparison

Total Papers: 5

Comparison

• {'title': 'Consensus and Cooperation in Networked Multi-Agent Systems', 'research_area': 'Multi-Agent Systems', 'research_problem': 'Not Available', 'methodology': ['Not Identified'], 'keywords': [], 'year': 2007, 'citation_count': 10391, 'quality_score': 50, 'quality_classification': 'Good'}

• {'title': 'Why Do Multi-Agent LLM Systems Fail?', 'research_area': 'Retrieval-Augmented Generation', 'research_problem': 'This gap highlights a critical need for a principled understanding of why MAS fail.', 'methodology': ['Framework', 'Experimental Evaluation', 'Dataset', 'Large Language Model', 'Retrieval-Augmented Generation'], 'keywords': ['multi-agent', 'mast-data', 'dataset', 'failure', 'agent', 'benchmark', 'inter-agent', 'inter-annotator', 'llm-as-a-judge', 'mast', 'annotation', 'systematic', 'agreement', 'better', 'design'], 'year': 2025, 'citation_count': 542, 'quality_score': 100, 'quality_classification': 'Excellent'}

• {'title': 'Which Agent Causes Task Failures and When? On Automated Failure Attribution of LLM Multi-Agent Systems', 'research_area': 'Multi-Agent Systems', 'research_problem': 'Failure attribution in LLM multi-agent systems-identifying the agent and step responsible for task failures-provides crucial clues for systems debugging but remains underexplored and labor-intensive.', 'methodology': ['Dataset', 'Large Language Model'], 'keywords': ['multi-agent', 'agent', 'failure', 'dataset', 'reasoning', 'attribution', 'failure-responsible', 'failures-provide', 'fine-grained', 'labor-intensive', 'systems-identifying', 'achieve', 'area', 'automated', 'steps'], 'year': 2025, 'citation_count': 147, 'quality_score': 90, 'quality_classification': 'Excellent'}

• {'title': 'Multi-Agent Actor-Critic for Mixed Cooperative-Competitive Environments', 'research_area': 'Reinforcement Learning', 'research_problem': 'We begin by analyzing the difficulty of traditional algorithms in the multi-agent case: Q-learning is challenged by an inherent non-stationarity of the environment, while policy gradient suffers from a variance that increases as the number of agents grows.', 'methodology': ['Algorithm', 'Reinforcement Learning'], 'keywords': ['multi-agent', 'agent', 'reinforcement', 'policy', 'actor-critic', 'non-stationarity', 'q-learning', 'coordination', 'able', 'adaptation', 'additionally', 'challenged', 'competitive', 'cooperative', 'difficulty'], 'year': 2017, 'citation_count': 6288, 'quality_score': 80, 'quality_classification': 'Very Good'}

• {'title': 'Red-Teaming LLM Multi-Agent Systems via Communication Attacks', 'research_area': 'Multi-Agent Systems', 'research_problem': 'Large Language Model-based Multi-Agent Systems (LLM-MAS) have revolutionized complex problem-solving capability by enabling sophisticated agent collaboration through message-based communications.', 'methodology': ['Framework', 'Experimental Evaluation', 'Large Language Model'], 'keywords': ['multi-agent', 'agent', 'communication', 'llm-ma', 'agent-in-the-middle', 'communication-based', 'contextually-aware', 'inter-agent', 'llm-powered', 'message-based', 'model-based', 'problem-solving', 'real-world', 'role-restricted', 'attack'], 'year': 2025, 'citation_count': 128, 'quality_score': 90, 'quality_classification': 'Excellent'}

Highest Cited Paper

title: Consensus and Cooperation in Networked Multi-Agent Systems

authors: ['R. Olfati-Saber', 'J. Fax', 'R. Murray']

year: 2007

citation_count: 10391

summary: Abstract not available for analysis.

research_problem: Not Available

methodology: ['Not Identified']

key_contributions: []

future_work: []

keywords: []

research_area: Multi-Agent Systems

paper_score: 50

paper_quality: Good

Latest Paper

title: Why Do Multi-Agent LLM Systems Fail?

authors: ['M. Cemri', 'Melissa Z. Pan', 'Shuyi Yang', 'Lakshya A. Agrawal', 'Bhavya Chopra', 'Rishabh Tiwari', 'Kurt Keutzer', 'Aditya G. Parameswaran', 'Dan Klein', 'K. Ramchandran', 'Matei A. Zaharia', 'Joseph E. Gonzalez', 'Ion Stoica']

year: 2025

citation_count: 542

summary: Despite enthusiasm for Multi-Agent LLM Systems (MAS), their performance gains on popular benchmarks are often minimal. This gap highlights a critical need for a principled understanding of why MAS fail. Addressing this question requires systematic identification and analysis of failure patterns. We introduce MAST-Data, a comprehensive dataset of 1600+ annotated traces collected across 7 popular MAS frameworks. MAST-Data is the first multi-agent system dataset to outline the failure dynamics in MAS for guiding the development of better future systems. To enable systematic classification of failures for MAST-Data, we build the first Multi-Agent System Failure Taxonomy (MAST). We develop MAST through rigorous analysis of 150 traces, guided closely by expert human annotators and validated by high inter-annotator agreement (kappa = 0.88). This process identifies 14 unique modes, clustered into 3 categories: (i) system design issues, (ii) inter-agent misalignment, and (iii) task verification. To enable scalable annotation, we develop an LLM-as-a-Judge pipeline with high agreement with human annotations. We leverage MAST and MAST-Data to analyze failure patterns across models (GPT4, Claude 3, Qwen2.5, CodeLlama) and tasks (coding, math, general agent), demonstrating improvement headrooms from better MAS design. Our analysis provides insights revealing that identified failures require more sophisticated solutions, highlighting a clear

research_problem: This gap highlights a critical need for a principled understanding of why MAS fail.

methodology: ['Framework', 'Experimental Evaluation', 'Dataset', 'Large Language Model', 'Retrieval-Augmented Generation']

key_contributions: ['We introduce MAST-Data, a comprehensive dataset of 1600+ annotated traces collected across 7 popular MAS frameworks.', 'We develop MAST through rigorous analysis of 150 traces, guided closely by expert human annotators and validated by high inter-annotator agreement (kappa = 0.88).', 'To enable scalable annotation, we develop an LLM-as-a-Judge pipeline with high agreement with human annotations.']

future_work: []

keywords: ['multi-agent', 'mast-data', 'dataset', 'failure', 'agent', 'benchmark', 'inter-agent', 'inter-annotator', 'llm-as-a-judge', 'mast', 'annotation', 'systematic', 'agreement', 'better', 'design']

research_area: Retrieval-Augmented Generation

paper_score: 100

paper_quality: Excellent

Common Methodologies

• Algorithm

• Dataset

• Experimental Evaluation

• Framework

• Large Language Model

• Not Identified

• Reinforcement Learning

• Retrieval-Augmented Generation

Research Areas

• Multi-Agent Systems

• Reinforcement Learning

• Retrieval-Augmented Generation

Common Keywords

• able

• achieve

• actor-critic

• adaptation

• additionally

• agent

• agent-in-the-middle

• agreement

• annotation

• area

• attack

• attribution

• automated

• benchmark

• better

• challenged

- communication
- communication-based
- competitive
- contextually-aware
- cooperative
- coordination
- dataset
- design
- difficulty
- failure
- failure-responsible
- failures-provide
- fine-grained
- inter-agent
- inter-annotator
- labor-intensive
- llm-as-a-judge
- llm-ma
- llm-powered
- mast
- mast-data
- message-based
- model-based
- multi-agent
- non-stationarity
- policy
- problem-solving
- q-learning
- real-world
- reasoning
- reinforcement
- role-restricted
- steps
- systematic
- systems-identifying

Methodology Frequency

Large Language Model: 3

Framework: 2

Experimental Evaluation: 2

Dataset: 2

Not Identified: 1

Retrieval-Augmented Generation: 1

Algorithm: 1

Reinforcement Learning: 1

Most Common Methodology: Large Language Model

Quality Distribution

Good: 1

Excellent: 3

Very Good: 1

Publication Years

2007: 1

2017: 1

2025: 3

Citation Statistics

highest: 10391

lowest: 128

average: 3499.2

total: 17496

Comparison Summary

total_papers: 5

most_common_methodology: Large Language Model

research_areas: ['Multi-Agent Systems', 'Reinforcement Learning', 'Retrieval-Augmented Generation']

quality_distribution: {'Good': 1, 'Excellent': 3, 'Very Good': 1}

citation_statistics: {'highest': 10391, 'lowest': 128, 'average': 3499.2, 'total': 17496}

highest_cited_title: Consensus and Cooperation in Networked Multi-Agent Systems

latest_paper_title: Why Do Multi-Agent LLM Systems Fail?

Research Gap Analysis

Total Papers: 5

Research Areas

- Multi-Agent Systems
- Reinforcement Learning
- Retrieval-Augmented Generation

Common Keywords

- agent
- multi-agent
- dataset
- failure
- inter-agent
- able
- achieve
- actor-critic
- adaptation
- additionally
- agent-in-the-middle
- agreement
- annotation
- area
- attack
- attribution
- automated
- benchmark
- better
- challenged
- communication
- communication-based
- competitive
- contextually-aware
- cooperative
- coordination
- design
- difficulty
- failure-responsible
- failures-provide
- fine-grained
- inter-annotator
- labor-intensive

- llm-as-a-judge
- llm-ma
- llm-powered
- mast
- mast-data
- message-based
- model-based
- non-stationarity
- policy
- problem-solving
- q-learning
- real-world
- reasoning
- reinforcement
- role-restricted
- steps
- systematic
- systems-identifying

Future Work

- These results highlight the task's complexity and the need for further research in this area.

Research Area Distribution

Multi-Agent Systems: 3

Retrieval-Augmented Generation: 1

Reinforcement Learning: 1

Keyword Frequency

agent: 4

multi-agent: 4

dataset: 2

failure: 2

inter-agent: 2

able: 1

achieve: 1

actor-critic: 1

adaptation: 1

additionally: 1

agent-in-the-middle: 1

agreement: 1
annotation: 1
area: 1
attack: 1
attribution: 1
automated: 1
benchmark: 1
better: 1
challenged: 1
communication: 1
communication-based: 1
competitive: 1
contextually-aware: 1
cooperative: 1
coordination: 1
design: 1
difficulty: 1
failure-responsible: 1
failures-provide: 1
fine-grained: 1
inter-annotator: 1
labor-intensive: 1
llm-as-a-judge: 1
llm-ma: 1
llm-powered: 1
mast: 1
mast-data: 1
message-based: 1
model-based: 1
non-stationarity: 1
policy: 1
problem-solving: 1
q-learning: 1
real-world: 1
reasoning: 1
reinforcement: 1
role-restricted: 1
steps: 1

systematic: 1

systems-identifying: 1

Research Trends

• agent

• multi-agent

• dataset

• failure

• inter-agent

• able

• achieve

• actor-critic

• adaptation

• additionally

• agent-in-the-middle

• agreement

• annotation

• area

• attack

Gap Categories

datasets: ['Despite enthusiasm for Multi-Agent LLM Systems (MAS), their performance gains on popular benchmarks are often minimal. This gap highlights a critical need for a principled understanding of why MAS fail. Addressing this question requires systematic identification and analysis of failure patterns. We introduce MAST-Data, a comprehensive dataset of 1600+ annotated traces collected across 7 popular MAS frameworks. MAST-Data is the first multi-agent system dataset to outline the failure dynamics in MAS for guiding the development of better future systems. To enable systematic classification of failures for MAST-Data, we build the first Multi-Agent System Failure Taxonomy (MAST). We develop MAST through rigorous analysis of 150 traces, guided closely by expert human annotators and validated by high inter-annotator agreement (kappa = 0.88). This process identifies 14 unique modes, clustered into 3 categories: (i) system design issues, (ii) inter-agent misalignment, and (iii) task verification. To enable scalable annotation, we develop an LLM-as-a-Judge pipeline with high agreement with human annotations. We leverage MAST and MAST-Data to analyze failure patterns across models (GPT4, Claude 3, Qwen2.5, CodeLlama) and tasks (coding, math, general agent), demonstrating improvement headrooms from better MAS design. Our analysis provides insights revealing that identified failures require more sophisticated solutions, highlighting a clear', 'Dataset', 'We introduce MAST-Data, a comprehensive dataset of 1600+ annotated traces collected across 7 popular MAS frameworks.', "Failure attribution in LLM multi-agent systems-identifying the agent and step responsible for task failures-provides crucial clues for systems debugging but remains underexplored and labor-intensive. In this paper, we propose and formulate a new research area: automated failure attribution for LLM multi-agent systems. To support this initiative, we introduce the Who&When; dataset, comprising extensive failure logs from 127 LLM multi-agent systems with fine-grained annotations linking failures to specific agents and decisive error steps. Using the Who&When;, we develop and evaluate three automated failure attribution methods, summarizing their corresponding pros and cons. The best method achieves 53.5% accuracy in identifying failure-responsible agents but only 14.2% in pinpointing failure steps, with some methods performing below random. Even

SOTA reasoning models, such as OpenAI o1 and DeepSeek R1, fail to achieve practical usability. These results highlight the task's complexity and the need for further research in this area. Code and dataset are available at https://github.com/mingyin1/Agents_Failure_Attribution", 'To support this initiative, we introduce the Who&When; dataset, comprising extensive failure logs from 127 LLM multi-agent systems with fine-grained annotations linking failures to specific agents and decisive error steps.']

models: ['Why Do Multi-Agent LLM Systems Fail?', 'Despite enthusiasm for Multi-Agent LLM Systems (MAS), their performance gains on popular benchmarks are often minimal. This gap highlights a critical need for a principled understanding of why MAS fail. Addressing this question requires systematic identification and analysis of failure patterns. We introduce MAST-Data, a comprehensive dataset of 1600+ annotated traces collected across 7 popular MAS frameworks. MAST-Data is the first multi-agent system dataset to outline the failure dynamics in MAS for guiding the development of better future systems. To enable systematic classification of failures for MAST-Data, we build the first Multi-Agent System Failure Taxonomy (MAST). We develop MAST through rigorous analysis of 150 traces, guided closely by expert human annotators and validated by high inter-annotator agreement (kappa = 0.88). This process identifies 14 unique modes, clustered into 3 categories: (i) system design issues, (ii) inter-agent misalignment, and (iii) task verification. To enable scalable annotation, we develop an LLM-as-a-Judge pipeline with high agreement with human annotations. We leverage MAST and MAST-Data to analyze failure patterns across models (GPT4, Claude 3, Qwen2.5, CodeLlama) and tasks (coding, math, general agent), demonstrating improvement headrooms from better MAS design. Our analysis provides insights revealing that identified failures require more sophisticated solutions, highlighting a clear', 'Large Language Model', 'To enable scalable annotation, we develop an LLM-as-a-Judge pipeline with high agreement with human annotations.', 'Which Agent Causes Task Failures and When? On Automated Failure Attribution of LLM Multi-Agent Systems', "Failure attribution in LLM multi-agent systems-identifying the agent and step responsible for task failures-provides crucial clues for systems debugging but remains underexplored and labor-intensive. In this paper, we propose and formulate a new research area: automated failure attribution for LLM multi-agent systems. To support this initiative, we introduce the Who&When; dataset, comprising extensive failure logs from 127 LLM multi-agent systems with fine-grained annotations linking failures to specific agents and decisive error steps. Using the Who&When;, we develop and evaluate three automated failure attribution methods, summarizing their corresponding pros and cons. The best method achieves 53.5% accuracy in identifying failure-responsible agents but only 14.2% in pinpointing failure steps, with some methods performing below random. Even SOTA reasoning models, such as OpenAI o1 and DeepSeek R1, fail to achieve practical usability. These results highlight the task's complexity and the need for further research in this area. Code and dataset are available at https://github.com/mingyin1/Agents_Failure_Attribution", 'Failure attribution in LLM multi-agent systems-identifying the agent and step responsible for task failures-provides crucial clues for systems debugging but remains underexplored and labor-intensive.', 'In this paper, we propose and formulate a new research area: automated failure attribution for LLM multi-agent systems.', 'To support this initiative, we introduce the Who&When; dataset, comprising extensive failure logs from 127 LLM multi-agent systems with fine-grained annotations linking failures to specific agents and decisive error steps.', 'Red-Teaming LLM Multi-Agent Systems via Communication Attacks', 'Large Language Model-based Multi-Agent Systems (LLM-MAS) have revolutionized complex problem-solving capability by enabling sophisticated agent collaboration through message-based communications. While the communication framework is crucial for agent coordination, it also introduces a critical yet unexplored security vulnerability. In this work, we introduce Agent-in-the-Middle (AiTM), a novel attack that exploits the fundamental communication mechanisms in LLM-MAS by intercepting and manipulating inter-agent messages. Unlike existing attacks that compromise individual agents, AiTM demonstrates how an adversary can compromise entire multi-agent systems by only manipulating the messages passing between agents. To enable the attack under the challenges of limited control and role-restricted communication format, we develop an LLM-powered adversarial agent with a reflection mechanism that generates contextually-aware malicious instructions. Our comprehensive evaluation across various frameworks, communication structures, and real-world applications demonstrates that LLM-MAS is vulnerable to communication-based attacks, highlighting the need for robust security measures in multi-agent systems.', 'Large Language Model-based Multi-Agent Systems (LLM-MAS) have revolutionized complex problem-solving capability by enabling sophisticated agent collaboration through message-based communications.', 'In this work, we introduce Agent-in-the-Middle (AiTM), a novel attack that exploits the fundamental communication mechanisms in LLM-MAS by intercepting and manipulating inter-agent messages.', 'To enable the attack under the challenges of limited control and role-restricted communication format, we develop an LLM-powered adversarial agent with a reflection mechanism that generates contextually-aware malicious instructions.']

retrieval: ['Retrieval-Augmented Generation']

grounding: []

hallucination: []

knowledge_freshness: []

evaluation: ['Despite enthusiasm for Multi-Agent LLM Systems (MAS), their performance gains on popular benchmarks are often minimal. This gap highlights a critical need for a principled understanding of why MAS fail. Addressing this question requires systematic identification and analysis of failure patterns. We introduce MAST-Data, a comprehensive dataset of 1600+ annotated traces collected across 7 popular MAS frameworks. MAST-Data is the first multi-agent system dataset to outline the failure dynamics in MAS for guiding the development of better future systems. To enable systematic classification of failures for MAST-Data, we build the first Multi-Agent System Failure Taxonomy (MAST). We develop MAST through rigorous analysis of 150 traces, guided closely by expert human annotators and validated by high inter-annotator agreement (kappa = 0.88). This process identifies 14 unique modes, clustered into 3 categories: (i) system design issues, (ii) inter-agent misalignment, and (iii) task verification. To enable scalable annotation, we develop an LLM-as-a-Judge pipeline with high agreement with human annotations. We leverage MAST and MAST-Data to analyze failure patterns across models (GPT4, Claude 3, Qwen2.5, CodeLlama) and tasks (coding, math, general agent), demonstrating improvement headrooms from better MAS design. Our analysis provides insights revealing that identified failures require more sophisticated solutions, highlighting a clear', 'Experimental Evaluation', "Failure attribution in LLM multi-agent systems-identifying the agent and step responsible for task failures-provides crucial clues for systems debugging but remains underexplored and labor-intensive. In this paper, we propose and formulate a new research area: automated failure attribution for LLM multi-agent systems. To support this initiative, we introduce the Who&When; dataset, comprising extensive failure logs from 127 LLM multi-agent systems with fine-grained annotations linking failures to specific agents and decisive error steps. Using the Who&When;, we develop and evaluate three automated failure attribution methods, summarizing their corresponding pros and cons. The best method achieves 53.5% accuracy in identifying failure-responsible agents but only 14.2% in pinpointing failure steps, with some methods performing below random. Even SOTA reasoning models, such as OpenAI o1 and DeepSeek R1, fail to achieve practical usability. These results highlight the task's complexity and the need for further research in this area. Code and dataset are available at https://github.com/mingyin1/Agents_Failure_Attribution", 'Using the Who&When;, we develop and evaluate three automated failure attribution methods, summarizing their corresponding pros and cons.', 'Large Language Model-based Multi-Agent Systems (LLM-MAS) have revolutionized complex problem-solving capability by enabling sophisticated agent collaboration through message-based communications. While the communication framework is crucial for agent coordination, it also introduces a critical yet unexplored security vulnerability. In this work, we introduce Agent-in-the-Middle (AiTM), a novel attack that exploits the fundamental communication mechanisms in LLM-MAS by intercepting and manipulating inter-agent messages. Unlike existing attacks that compromise individual agents, AiTM demonstrates how an adversary can compromise entire multi-agent systems by only manipulating the messages passing between agents. To enable the attack under the challenges of limited control and role-restricted communication format, we develop an LLM-powered adversarial agent with a reflection mechanism that generates contextually-aware malicious instructions. Our comprehensive evaluation across various frameworks, communication structures, and real-world applications demonstrates that LLM-MAS is vulnerable to communication-based attacks, highlighting the need for robust security measures in multi-agent systems.']

applications: ['Large Language Model-based Multi-Agent Systems (LLM-MAS) have revolutionized complex problem-solving capability by enabling sophisticated agent collaboration through message-based communications. While the communication framework is crucial for agent coordination, it also introduces a critical yet unexplored security vulnerability. In this work, we introduce Agent-in-the-Middle (AiTM), a novel attack that exploits the fundamental communication mechanisms in LLM-MAS by intercepting and manipulating inter-agent messages. Unlike existing attacks that compromise individual agents, AiTM demonstrates how an adversary can compromise entire multi-agent systems by only manipulating the messages passing between agents. To enable the attack under the challenges of limited control and role-restricted communication format, we develop an LLM-powered adversarial agent with a reflection mechanism that generates contextually-aware malicious instructions. Our comprehensive evaluation across various frameworks, communication structures, and real-world applications demonstrates that LLM-MAS is vulnerable to communication-based attacks, highlighting the need for robust security measures in multi-agent systems.']

coordination: ['Consensus and Cooperation in Networked Multi-Agent Systems', 'Why Do Multi-Agent LLM Systems Fail?', 'Despite enthusiasm for Multi-Agent LLM Systems (MAS), their performance gains on popular benchmarks are often minimal. This gap highlights a critical need for a principled understanding of why MAS fail. Addressing this question requires systematic identification and analysis of failure patterns. We introduce MAST-Data, a comprehensive dataset of 1600+ annotated traces collected across 7 popular MAS frameworks. MAST-Data is the first multi-agent system dataset to outline the failure dynamics in MAS for guiding the development of better future systems. To enable systematic classification of failures for MAST-Data, we build the first Multi-Agent System Failure Taxonomy (MAST). We develop MAST through rigorous analysis of 150 traces, guided closely by expert human annotators and validated by high inter-annotator agreement (kappa = 0.88). This process identifies 14 unique modes, clustered into 3 categories: (i) system design issues, (ii) inter-agent misalignment, and (iii) task verification. To enable scalable annotation, we develop an LLM-as-a-Judge pipeline with high agreement with human annotations. We leverage MAST and MAST-Data to analyze failure patterns across models (GPT4, Claude 3, Qwen2.5, CodeLlama) and tasks (coding, math, general agent), demonstrating improvement headrooms from better MAS design. Our analysis provides insights revealing that identified failures require more sophisticated solutions, highlighting a clear', 'Which Agent Causes Task Failures and When? On Automated Failure Attribution of LLM Multi-Agent Systems', "Failure attribution in LLM multi-agent systems-identifying the agent and step responsible for task failures-provides crucial clues for systems debugging but remains underexplored and labor-intensive. In this paper, we propose and formulate a new research area: automated failure attribution for LLM multi-agent systems. To support this initiative, we introduce the Who&When; dataset, comprising extensive failure logs from 127 LLM multi-agent systems with fine-grained annotations linking failures to specific agents and decisive error steps. Using the Who&When;, we develop and evaluate three automated failure attribution methods, summarizing their corresponding pros and cons. The best method achieves 53.5% accuracy in identifying failure-responsible agents but only 14.2% in pinpointing failure steps, with some methods performing below random. Even SOTA reasoning models, such as OpenAI o1 and DeepSeek R1, fail to achieve practical usability. These results highlight the task's complexity and the need for further research in this area. Code and dataset are available at https://github.com/mingyin1/Agents_Failure_Attribution", 'Failure attribution in LLM multi-agent systems-identifying the agent and step responsible for task failures-provides crucial clues for systems debugging but remains underexplored and labor-intensive.', 'In this paper, we propose and formulate a new research area: automated failure attribution for LLM multi-agent systems.', 'To support this initiative, we introduce the Who&When; dataset, comprising extensive failure logs from 127 LLM multi-agent systems with fine-grained annotations linking failures to specific agents and decisive error steps.', 'Multi-Agent Actor-Critic for Mixed Cooperative-Competitive Environments', 'We explore deep reinforcement learning methods for multi-agent domains. We begin by analyzing the difficulty of traditional algorithms in the multi-agent case: Q-learning is challenged by an inherent non-stationarity of the environment, while policy gradient suffers from a variance that increases as the number of agents grows. We then present an adaptation of actor-critic methods that considers action policies of other agents and is able to successfully learn policies that require complex multi-agent coordination. Additionally, we introduce a training regimen utilizing an ensemble of policies for each agent that leads to more robust multi-agent policies. We show the strength of our approach compared to existing methods in cooperative as well as competitive scenarios, where agent populations are able to discover various physical and informational coordination strategies.', 'We begin by analyzing the difficulty of traditional algorithms in the multi-agent case: Q-learning is challenged by an inherent non-stationarity of the environment, while policy gradient suffers from a variance that increases as the number of agents grows.', 'Additionally, we introduce a training regimen utilizing an ensemble of policies for each agent that leads to more robust multi-agent policies.', 'We show the strength of our approach compared to existing methods in cooperative as well as competitive scenarios, where agent populations are able to discover various physical and informational coordination strategies.', 'Red-Teaming LLM Multi-Agent Systems via Communication Attacks', 'Large Language Model-based Multi-Agent Systems (LLM-MAS) have revolutionized complex problem-solving capability by enabling sophisticated agent collaboration through message-based communications. While the communication framework is crucial for agent coordination, it also introduces a critical yet unexplored security vulnerability. In this work, we introduce Agent-in-the-Middle (AiTM), a novel attack that exploits the fundamental communication mechanisms in LLM-MAS by intercepting and manipulating inter-agent messages. Unlike existing attacks that compromise individual agents, AiTM demonstrates how an adversary can compromise entire multi-agent systems by only manipulating the messages passing between agents. To enable the attack under the challenges of limited control and role-restricted communication format, we develop an LLM-powered adversarial agent with a reflection mechanism that generates contextually-aware malicious instructions. Our comprehensive evaluation across various frameworks, communication structures, and real-world applications demonstrates that LLM-MAS is vulnerable to communication-based attacks, highlighting the need for robust security measures in multi-agent systems.', 'Large Language Model-based Multi-Agent Systems (LLM-MAS) have revolutionized complex problem-solving capability by enabling sophisticated agent collaboration through message-based communications.', 'In this work, we introduce Agent-in-the-Middle (AiTM), a novel attack that exploits the fundamental communication mechanisms in LLM-MAS by intercepting and manipulating inter-agent messages.', 'Unlike existing attacks that compromise individual agents, AiTM demonstrates how an adversary can compromise entire multi-agent systems by only manipulating the messages passing between agents.', 'To enable the attack under the challenges of limited control and role-restricted communication format, we develop an LLM-powered adversarial agent with a reflection mechanism that generates contextually-aware malicious instructions.']

security: ['Red-Teaming LLM Multi-Agent Systems via Communication Attacks', 'Large Language Model-based Multi-Agent Systems (LLM-MAS) have revolutionized complex problem-solving capability by enabling sophisticated agent collaboration through message-based communications. While the communication framework is crucial for agent coordination, it also introduces a critical yet unexplored security vulnerability. In this work, we introduce Agent-in-the-Middle (AiTM), a novel attack that exploits the fundamental communication mechanisms in LLM-MAS by intercepting and manipulating inter-agent messages. Unlike existing attacks that compromise individual agents, AiTM demonstrates how an adversary can compromise entire multi-agent systems by only manipulating the messages passing between agents. To enable the attack under the challenges of limited control and role-restricted communication format, we develop an LLM-powered adversarial agent with a reflection mechanism that generates contextually-aware malicious instructions. Our comprehensive evaluation across various frameworks, communication structures, and real-world applications demonstrates that LLM-MAS is vulnerable to communication-based attacks, highlighting the need for robust security measures in multi-agent systems.', 'In this work, we introduce Agent-in-the-Middle (AiTM), a novel attack that exploits the fundamental communication mechanisms in LLM-MAS by intercepting and manipulating inter-agent messages.', 'Unlike existing attacks that compromise individual agents, AiTM demonstrates how an adversary can compromise entire multi-agent systems by only manipulating the messages passing between agents.', 'To enable the attack under the challenges of limited control and role-restricted communication format, we develop an LLM-powered adversarial agent with a reflection mechanism that generates contextually-aware malicious instructions.']

scalability: ['Despite enthusiasm for Multi-Agent LLM Systems (MAS), their performance gains on popular benchmarks are often minimal. This gap highlights a critical need for a principled understanding of why MAS fail. Addressing this question requires systematic identification and analysis of failure patterns. We introduce MAST-Data, a comprehensive dataset of 1600+ annotated traces collected across 7 popular MAS frameworks. MAST-Data is the first multi-agent system dataset to outline the failure dynamics in MAS for guiding the development of better future systems. To enable systematic classification of failures for MAST-Data, we build the first Multi-Agent System Failure Taxonomy (MAST). We develop MAST through rigorous analysis of 150 traces, guided closely by expert human annotators and validated by high inter-annotator agreement (kappa = 0.88). This process identifies 14 unique modes, clustered into 3 categories: (i) system design issues, (ii) inter-agent misalignment, and (iii) task verification. To enable scalable annotation, we develop an LLM-as-a-Judge pipeline with high agreement with human annotations. We leverage MAST and MAST-Data to analyze failure patterns across models (GPT4, Claude 3, Qwen2.5, CodeLlama) and tasks (coding, math, general agent), demonstrating improvement headrooms from better MAS design. Our analysis provides insights revealing that identified failures require more sophisticated solutions, highlighting a clear', 'To enable scalable annotation, we develop an LLM-as-a-Judge pipeline with high agreement with human annotations.']

explainability: []

reasoning: []

memory: []

planning: []

general: ["These results highlight the task's complexity and the need for further research in this area."]

Research Gaps

• Retrieval quality and the ability to consistently retrieve relevant information for complex queries remain important research challenges.

• More robust and standardized evaluation and benchmarking methods are required.

• Larger, more diverse, and domain-specific datasets are needed for reliable development and evaluation.

• Current models and architectures still require improvements in reliability, adaptability, and generalization.

• Reliable coordination and communication between multiple agents remains an open research challenge.

• Security, privacy, trust, and robustness remain important unresolved challenges.

• Scalability, computational cost, and efficient resource usage remain challenges for practical deployment.

• Further validation in real-world applications and deployment environments is required.

• Despite enthusiasm for Multi-Agent LLM Systems (MAS), their performance gains on popular benchmarks are often minimal.

• This process identifies 14 unique modes, clustered into 3 categories: (i) system design issues, (ii) inter-agent misalignment, and (iii) task verification.

Emerging Topics

• agent

• multi-agent

• dataset

• failure

• inter-agent

Recommendations

• Explore more reliable autonomous and multi-agent systems.

Summary

dominant_area: Multi-Agent Systems

top_trend: agent

future_work_items: 1

research_gap_count: 10

Experiment Plan

Research Areas

• Multi-Agent Systems

• Reinforcement Learning

• Retrieval-Augmented Generation

Recommended Datasets

• AgentBench

• BEIR

• GAIA Benchmark

• HotpotQA

• MS MARCO

• Multi-Agent Benchmark

• Natural Questions

• OpenAI Evals

• TriviaQA

Baseline Models

• AutoGen

• BM25 Retrieval

• CrewAI

• Dense Retrieval RAG

• Hybrid Search RAG

• LangGraph

• Naive RAG

• ReAct

• Single-Agent LLM

- Vanilla RAG

Evaluation Metrics

- Accuracy
- Agent Coordination Score
- Answer Relevance
- Context Relevance
- F1-Score
- Faithfulness
- Groundedness
- Planning Accuracy
- Precision
- ROC-AUC
- Recall
- Retrieval Precision
- Retrieval Recall
- Task Success Rate
- Tool-Use Accuracy

Hardware Requirements

cpu: Intel Core i7 / AMD Ryzen 7

ram: 16 GB

gpu: NVIDIA RTX 3060 or better

storage: 100 GB SSD

Validation Strategy

- Train-Test Split
- K-Fold Cross Validation
- Hyperparameter Tuning
- Statistical Significance Testing

Experimental Workflow

- Collect Knowledge Documents
- Preprocess and Chunk Documents
- Build Document Index
- Configure Retrieval System
- Retrieve Relevant Context
- Generate Grounded Response
- Evaluate Retrieval Quality

• Evaluate Answer Quality
• Compare Baseline and Proposed RAG
• Analyze Errors
• Draw Conclusions

Report Summary

Dominant Research Area: Multi-Agent Systems

Top Research Trend: agent

Highest Cited Paper: Consensus and Cooperation in Networked Multi-Agent Systems

Latest Paper: Why Do Multi-Agent LLM Systems Fail?

Recommended Datasets

• AgentBench
• BEIR
• GAIA Benchmark
• HotpotQA
• MS MARCO
• Multi-Agent Benchmark
• Natural Questions
• OpenAI Evals
• TriviaQA

Baseline Models

• AutoGen
• BM25 Retrieval
• CrewAI
• Dense Retrieval RAG
• Hybrid Search RAG
• LangGraph
• Naive RAG
• ReAct
• Single-Agent LLM
• Vanilla RAG

Evaluation Metrics

• Accuracy
• Agent Coordination Score
• Answer Relevance

- Context Relevance
- F1-Score
- Faithfulness
- Groundedness
- Planning Accuracy
- Precision
- ROC-AUC
- Recall
- Retrieval Precision
- Retrieval Recall
- Task Success Rate
- Tool-Use Accuracy

Conclusion

The report has been generated using a personalized multi-agent pipeline. The explanation style was adapted according to the Intermediate profile to improve readability and usefulness.