Research Analysis Report

Research Topic: RAG

Total Papers: 5

Generated By: Report Generation Agent

Generated At: 2026-08-31 21:31:32

Executive Summary

This report provides a technical overview of 'RAG' including methodologies, research gaps and experiment planning.

Literature Review

Introduction: This literature review summarizes findings from 5 research papers related to the selected topic. It highlights research trends, methodologies, key contributions, future directions, and overall observations.

Paper Summaries

• {'title': 'From Local to Global: A Graph RAG Approach to Query-Focused Summarization', 'year': 2024, 'summary': 'The use of retrieval-augmented generation (RAG) to retrieve relevant information from an external knowledge source enables large language models (LLMs) to answer questions over private and/or previously unseen document collections. However, RAG fails on global questions directed at an entire text corpus, such as"What are the main themes in the dataset?", since this is inherently a query-focused summarization (QFS) task, rather than an explicit retrieval task. Prior QFS methods, meanwhile, do not scale to the quantities of text indexed by typical RAG systems. To combine the strengths of these contrasting methods, we propose GraphRAG, a graph-based approach to question answering over private text corpora that scales with both the generality of user questions and the quantity of source text. Our approach uses an LLM to build a graph index in two stages: first, to derive an entity knowledge graph from the source documents, then to pregenerate community summaries for all groups of closely related entities. Given a question, each community summary is used to generate a partial response, before all partial responses are again summarized in a final response to the user. For a class of global sensemaking questions over datasets in the 1 million token range, we show that'}

• {'title': 'Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection', 'year': 2023, 'summary': "Despite their remarkable capabilities, large language models (LLMs) often produce responses containing factual inaccuracies due to their sole reliance on the parametric knowledge they encapsulate. Retrieval-Augmented Generation (RAG), an ad hoc approach that augments LMs with retrieval of relevant knowledge, decreases such issues. However, indiscriminately retrieving and incorporating a fixed number of retrieved passages, regardless of whether retrieval is necessary, or passages are relevant, diminishes LM versatility or can lead to unhelpful response generation. We introduce a new framework called Self-Reflective Retrieval-Augmented Generation (Self-RAG) that enhances an LM's quality and factuality through retrieval and self-reflection. Our framework trains a single arbitrary LM that adaptively retrieves passages on-demand, and generates and reflects on retrieved passages and its own generations using special tokens, called reflection tokens. Generating reflection tokens makes the LM controllable during the inference phase, enabling it to tailor its behavior to diverse task requirements. Experiments show that Self-RAG (7B and 13B parameters) significantly outperforms state-of-the-art LLMs and retrieval-augmented models on a diverse set of tasks. Specifically, Self-RAG outperforms ChatGPT and retrieval-augmented Llama2-chat on Open-domain QA, reasoning and fact verification tasks, and it shows significant gains in improving factuality and citation accuracy for long-form generations relative to these models."}

- {'title': 'From RAG to Memory: Non-Parametric Continual Learning for Large Language Models', 'year': 2025, 'summary': 'Our ability to continuously acquire, organize, and leverage knowledge is a key feature of human intelligence that AI systems must approximate to unlock their full potential. Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information. However, its reliance on vector retrieval hinders its ability to mimic the dynamic and interconnected nature of human long-term memory. Recent RAG approaches augment vector embeddings with various structures like knowledge graphs to address some of these gaps, namely sense-making and associativity. However, their performance on more basic factual memory tasks drops considerably below standard RAG. We address this unintended deterioration and propose HippoRAG 2, a framework that outperforms standard RAG comprehensively on factual, sense-making, and associative memory tasks. HippoRAG 2 builds upon the Personalized PageRank algorithm used in HippoRAG and enhances it with deeper passage integration and more effective online use of an LLM. This combination pushes this RAG system closer to the effectiveness of human long-term memory, achieving a 7% improvement in associative memory tasks over the state-of-the-art embedding model while also exhibiting superior factual knowledge and sense-making memory capabilities. This work paves the way for non-parametric continual learning'}

- {'title': 'Agentic Retrieval-Augmented Generation: A Survey on Agentic RAG', 'year': 2025, 'summary': 'Large Language Models (LLMs) have advanced artificial intelligence by enabling human-like text generation and natural language understanding. However, their reliance on static training data limits their ability to respond to dynamic, real-time queries, resulting in outdated or inaccurate outputs. Retrieval-Augmented Generation (RAG) has emerged as a solution, enhancing LLMs by integrating real-time data retrieval to provide contextually relevant and up-to-date responses. Despite its promise, traditional RAG systems are constrained by static workflows and lack the adaptability required for multi-step reasoning and complex task management. Agentic Retrieval-Augmented Generation (Agentic RAG) transcends these limitations by embedding autonomous AI agents into the RAG pipeline. These agents leverage agentic design patterns reflection, planning, tool use, and multi-agent collaboration to dynamically manage retrieval strategies, iteratively refine contextual understanding, and adapt workflows through operational structures ranging from sequential steps to adaptive collaboration. This integration enables Agentic RAG systems to deliver flexibility, scalability, and context-awareness across diverse applications. This paper presents an analytical survey of Agentic RAG systems. It traces the evolution of RAG paradigms, introduces a principled taxonomy of Agentic RAG architectures based on agent cardinality, control structure, autonomy, and knowledge representation, and provides a comparative analysis of design trade-offs across existing frameworks. The survey examines'}

- {'title': 'Retrieval-Augmented Generation (RAG) Chatbots for Education: A Survey of Applications', 'year': 2025, 'summary': 'Retrieval-Augmented Generation (RAG) overcomes the main barrier for the adoption of LLM-based chatbots in education: hallucinations. The uncomplicated architecture of RAG chatbots makes it relatively easy to implement chatbots that serve specific purposes and thus are capable of addressing various needs in the educational domain. With five years having passed since the introduction of RAG, the time has come to check the progress attained in its adoption in education. This paper identifies 47 papers dedicated to RAG chatbots' uses for various kinds of educational purposes, which are analyzed in terms of their character, the target of the support provided by the chatbots, the thematic scope of the knowledge accessible via the chatbots, the underlying large language model, and the character of their evaluation.'}

Research Trends

- accessible
- address

- addressing
- adoption
- agent
- agentic
- architecture
- associative
- autonomous
- character
- chatbot
- community
- context-awareness
- continual
- dataset
- document
- education
- educational
- entity
- evaluation
- factual
- factuality
- generation
- global
- graph
- graph-based
- hallucination
- hipporag
- human
- human-like
- introduction
- knowledge
- llm-based
- long-form
- long-term
- memory
- multi-agent
- multi-step
- non-parametric
- on-demand

- open-domain
- partial
- passage
- planning
- query-focused
- question
- real-time
- reasoning
- response
- retrieval
- retrieval-augmented
- self-rag
- self-reflection
- self-reflective
- sense-making
- source
- state-of-the-art
- tasks
- text
- token
- trade-off
- up-to-date

Research Gaps

None

Future Scope

None

Conclusion: Overall, the reviewed studies demonstrate significant progress in this research area while highlighting several opportunities for future work and further experimentation.

Methodology Comparison

Total Papers: 5

Comparison

• {'title': 'From Local to Global: A Graph RAG Approach to Query-Focused Summarization', 'research_area': 'Retrieval-Augmented Generation', 'research_problem': 'However, RAG fails on global questions directed at an entire text corpus, such as"What are the main themes in the dataset?", since this is inherently a query-focused summarization (QFS) task, rather than an explicit retrieval task.', 'methodology': ['Dataset', 'Large Language Model', 'Retrieval-Augmented

Generation'], 'keywords': ['dataset', 'graph', 'knowledge', 'question', 'graph-based', 'query-focused', 'retrieval-augmented', 'text', 'response', 'source', 'community', 'document', 'entity', 'global', 'partial'], 'year': 2024, 'citation_count': 2034, 'quality_score': 90, 'quality_classification': 'Excellent'}

- {'title': 'Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection', 'research_area': 'Retrieval-Augmented Generation', 'research_problem': 'Retrieval-Augmented Generation (RAG), an ad hoc approach that augments LMs with retrieval of relevant knowledge, decreases such issues.', 'methodology': ['Framework', 'Experimental Evaluation', 'Large Language Model', 'Retrieval-Augmented Generation'], 'keywords': ['retrieval-augmented', 'knowledge', 'generation', 'self-rag', 'reasoning', 'open-domain', 'self-reflection', 'self-reflective', 'state-of-the-art', 'passage', 'long-form', 'on-demand', 'retrieval', 'token', 'factuality'], 'year': 2023, 'citation_count': 2484, 'quality_score': 90, 'quality_classification': 'Excellent'}
- {'title': 'From RAG to Memory: Non-Parametric Continual Learning for Large Language Models', 'research_area': 'Retrieval-Augmented Generation', 'research_problem': 'Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information.', 'methodology': ['Framework', 'Algorithm', 'Large Language Model', 'Retrieval-Augmented Generation'], 'keywords': ['memory', 'knowledge', 'sense-making', 'graph', 'long-term', 'non-parametric', 'retrieval-augmented', 'state-of-the-art', 'factual', 'hipporag', 'human', 'tasks', 'associative', 'address', 'continual'], 'year': 2025, 'citation_count': 218, 'quality_score': 90, 'quality_classification': 'Excellent'}
- {'title': 'Agentic Retrieval-Augmented Generation: A Survey on Agentic RAG', 'research_area': 'Retrieval-Augmented Generation', 'research_problem': 'However, their reliance on static training data limits their ability to respond to dynamic, real-time queries, resulting in outdated or inaccurate outputs.', 'methodology': ['Survey', 'Framework', 'Architecture', 'Large Language Model', 'Retrieval-Augmented Generation'], 'keywords': ['multi-agent', 'agent', 'autonomous', 'agentic', 'retrieval-augmented', 'knowledge', 'planning', 'reasoning', 'real-time', 'context-awareness', 'human-like', 'multi-step', 'up-to-date', 'generation', 'trade-off'], 'year': 2025, 'citation_count': 405, 'quality_score': 90, 'quality_classification': 'Excellent'}
- {'title': 'Retrieval-Augmented Generation (RAG) Chatbots for Education: A Survey of Applications', 'research_area': 'Retrieval-Augmented Generation', 'research_problem': 'Retrieval-Augmented Generation (RAG) overcomes the main barrier for the adoption of LLM-based chatbots in education: hallucinations.', 'methodology': ['Architecture', 'Experimental Evaluation', 'Large Language Model', 'Retrieval-Augmented Generation'], 'keywords': ['chatbot', 'knowledge', 'retrieval-augmented', 'llm-based', 'educational', 'adoption', 'character', 'education', 'accessible', 'addressing', 'architecture', 'evaluation', 'generation', 'hallucination', 'introduction'], 'year': 2025, 'citation_count': 115, 'quality_score': 90, 'quality_classification': 'Excellent'}

Highest Cited Paper

title: Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection

authors: ['Akari Asai', 'Zeqiu Wu', 'Yizhong Wang', 'Avirup Sil', 'Hannaneh Hajishirzi']

year: 2023

citation_count: 2484

summary: Despite their remarkable capabilities, large language models (LLMs) often produce responses containing factual inaccuracies due to their sole reliance on the parametric knowledge they encapsulate. Retrieval-Augmented Generation (RAG), an ad hoc approach that augments LMs with retrieval of relevant knowledge, decreases such issues. However, indiscriminately retrieving and incorporating a fixed number of retrieved passages, regardless of whether retrieval is necessary, or passages are relevant, diminishes LM versatility or can lead to unhelpful response generation. We introduce a new framework called Self-Reflective Retrieval-Augmented Generation (Self-RAG) that enhances an LM's quality and factuality through retrieval and self-reflection. Our framework trains a single arbitrary LM that adaptively retrieves passages on-demand, and generates and reflects on retrieved passages and its own generations using special tokens, called reflection tokens. Generating reflection tokens makes the LM controllable during the inference phase, enabling it to tailor its behavior to diverse task requirements. Experiments show that Self-RAG (7B and 13B parameters) significantly outperforms state-of-the-art LLMs and retrieval-augmented models on a diverse set of tasks. Specifically, Self-RAG outperforms ChatGPT and retrieval-augmented Llama2-chat on Open-domain QA, reasoning and fact verification tasks, and it shows significant gains in improving factuality and citation accuracy for long-form generations relative to these models.

research_problem: Retrieval-Augmented Generation (RAG), an ad hoc approach that augments LMs with retrieval of relevant knowledge, decreases such issues.

methodology: ['Framework', 'Experimental Evaluation', 'Large Language Model', 'Retrieval-Augmented Generation']

key_contributions: ["We introduce a new framework called Self-Reflective Retrieval-Augmented Generation (Self-RAG) that enhances an LM's quality and factuality through retrieval and self-reflection.", 'Our framework trains a single arbitrary LM that adaptively retrieves passages on-demand, and generates and reflects on retrieved passages and its own generations using special tokens, called reflection tokens.', 'Experiments show that Self-RAG (7B and 13B parameters) significantly outperforms state-of-the-art LLMs and retrieval-augmented models on a diverse set of tasks.']

future_work: []

keywords: ['retrieval-augmented', 'knowledge', 'generation', 'self-rag', 'reasoning', 'open-domain', 'self-reflection', 'self-reflective', 'state-of-the-art', 'passage', 'long-form', 'on-demand', 'retrieval', 'token', 'factuality']

research_area: Retrieval-Augmented Generation

paper_score: 90

paper_quality: Excellent

Latest Paper

title: From RAG to Memory: Non-Parametric Continual Learning for Large Language Models

authors: ['Bernal Jiménez Gutiérrez', 'Yiheng Shu', 'Weijian Qi', 'Sizhe Zhou', 'Yu Su']

year: 2025

citation_count: 218

summary: Our ability to continuously acquire, organize, and leverage knowledge is a key feature of human intelligence that AI systems must approximate to unlock their full potential. Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information. However, its reliance on vector retrieval hinders its ability to mimic the dynamic and interconnected nature of human long-term memory. Recent RAG approaches augment vector embeddings with various structures like knowledge graphs to address some of these gaps, namely sense-making and associativity. However, their performance on more basic factual memory tasks drops considerably below standard RAG. We address this unintended deterioration and propose HippoRAG 2, a framework that outperforms standard RAG comprehensively on factual, sense-making, and associative memory tasks. HippoRAG 2 builds upon the Personalized PageRank algorithm used in HippoRAG and enhances it with deeper passage integration and more effective online use of an LLM. This combination pushes this RAG system closer to the effectiveness of human long-term memory, achieving a 7% improvement in associative memory tasks over the state-of-the-art embedding model while also exhibiting superior factual knowledge and sense-making memory capabilities. This work paves the way for non-parametric continual learning

research_problem: Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information.

methodology: ['Framework', 'Algorithm', 'Large Language Model', 'Retrieval-Augmented Generation']

key_contributions: ['We address this unintended deterioration and propose HippoRAG 2, a framework that outperforms standard RAG comprehensively on factual, sense-making, and associative memory tasks.']

future_work: []

keywords: ['memory', 'knowledge', 'sense-making', 'graph', 'long-term', 'non-parametric', 'retrieval-augmented', 'state-of-the-art', 'factual', 'hipporag', 'human', 'tasks', 'associative', 'address', 'continual']

research_area: Retrieval-Augmented Generation

paper_score: 90

paper_quality: Excellent

Common Methodologies

• Algorithm

• Architecture

• Dataset

• Experimental Evaluation

• Framework

• Large Language Model

• Retrieval-Augmented Generation

• Survey

Research Areas

• Retrieval-Augmented Generation

Common Keywords

• accessible

• address

• addressing

• adoption

• agent

• agentic

• architecture

• associative

• autonomous

• character

• chatbot

• community

• context-awareness

• continual

- dataset
- document
- education
- educational
- entity
- evaluation
- factual
- factuality
- generation
- global
- graph
- graph-based
- hallucination
- hipporag
- human
- human-like
- introduction
- knowledge
- llm-based
- long-form
- long-term
- memory
- multi-agent
- multi-step
- non-parametric
- on-demand
- open-domain
- partial
- passage
- planning
- query-focused
- question
- real-time
- reasoning
- response
- retrieval
- retrieval-augmented
- self-rag

- self-reflection
- self-reflective
- sense-making
- source
- state-of-the-art
- tasks
- text
- token
- trade-off
- up-to-date

Methodology Frequency

Large Language Model: 5

Retrieval-Augmented Generation: 5

Framework: 3

Experimental Evaluation: 2

Architecture: 2

Dataset: 1

Algorithm: 1

Survey: 1

Most Common Methodology: Large Language Model

Quality Distribution

Excellent: 5

Publication Years

2023: 1

2024: 1

2025: 3

Citation Statistics

highest: 2484

lowest: 115

average: 1051.2

total: 5256

Comparison Summary

total_papers: 5

most_common_methodology: Large Language Model

research_areas: ['Retrieval-Augmented Generation'] quality_distribution: {'Excellent': 5}

citation_statistics: {'highest': 2484, 'lowest': 115, 'average': 1051.2, 'total': 5256}

highest_cited_title: Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection

latest_paper_title: From RAG to Memory: Non-Parametric Continual Learning for Large Language Models

Research Gap Analysis

Total Papers: 5

Research Areas

• Retrieval-Augmented Generation

Common Keywords

• knowledge

• retrieval-augmented

• generation

• graph

• reasoning

• state-of-the-art

• accessible

• address

• addressing

• adoption

• agent

• agentic

• architecture

• associative

• autonomous

• character

• chatbot

• community

• context-awareness

• continual

• dataset

• document

• education

• educational

• entity

- evaluation
- factual
- factuality
- global
- graph-based
- hallucination
- hipporag
- human
- human-like
- introduction
- llm-based
- long-form
- long-term
- memory
- multi-agent
- multi-step
- non-parametric
- on-demand
- open-domain
- partial
- passage
- planning
- query-focused
- question
- real-time
- response
- retrieval
- self-rag
- self-reflection
- self-reflective
- sense-making
- source
- tasks
- text
- token
- trade-off
- up-to-date

Future Work

None

Research Area Distribution

Retrieval-Augmented Generation: 5

Keyword Frequency

knowledge: 5

retrieval-augmented: 5

generation: 3

graph: 2

reasoning: 2

state-of-the-art: 2

accessible: 1

address: 1

addressing: 1

adoption: 1

agent: 1

agentic: 1

architecture: 1

associative: 1

autonomous: 1

character: 1

chatbot: 1

community: 1

context-awareness: 1

continual: 1

dataset: 1

document: 1

education: 1

educational: 1

entity: 1

evaluation: 1

factual: 1

factuality: 1

global: 1

graph-based: 1

hallucination: 1

hipporag: 1

human: 1

human-like: 1

introduction: 1

llm-based: 1

long-form: 1

long-term: 1

memory: 1

multi-agent: 1

multi-step: 1

non-parametric: 1

on-demand: 1

open-domain: 1

partial: 1

passage: 1

planning: 1

query-focused: 1

question: 1

real-time: 1

response: 1

retrieval: 1

self-rag: 1

self-reflection: 1

self-reflective: 1

sense-making: 1

source: 1

tasks: 1

text: 1

token: 1

trade-off: 1

up-to-date: 1

Research Trends

• knowledge

• retrieval-augmented

• generation

• graph

• reasoning

• state-of-the-art
• accessible
• address
• addressing
• adoption
• agent
• agentic
• architecture
• associative
• autonomous

Gap Categories

datasets: ['The use of retrieval-augmented generation (RAG) to retrieve relevant information from an external knowledge source enables large language models (LLMs) to answer questions over private and/or previously unseen document collections. However, RAG fails on global questions directed at an entire text corpus, such as"What are the main themes in the dataset?", since this is inherently a query-focused summarization (QFS) task, rather than an explicit retrieval task. Prior QFS methods, meanwhile, do not scale to the quantities of text indexed by typical RAG systems. To combine the strengths of these contrasting methods, we propose GraphRAG, a graph-based approach to question answering over private text corpora that scales with both the generality of user questions and the quantity of source text. Our approach uses an LLM to build a graph index in two stages: first, to derive an entity knowledge graph from the source documents, then to pregenerate community summaries for all groups of closely related entities. Given a question, each community summary is used to generate a partial response, before all partial responses are again summarized in a final response to the user. For a class of global sensemaking questions over datasets in the 1 million token range, we show that', 'However, RAG fails on global questions directed at an entire text corpus, such as"What are the main themes in the dataset?", since this is inherently a query-focused summarization (QFS) task, rather than an explicit retrieval task.', 'Dataset']

models: ['The use of retrieval-augmented generation (RAG) to retrieve relevant information from an external knowledge source enables large language models (LLMs) to answer questions over private and/or previously unseen document collections. However, RAG fails on global questions directed at an entire text corpus, such as"What are the main themes in the dataset?", since this is inherently a query-focused summarization (QFS) task, rather than an explicit retrieval task. Prior QFS methods, meanwhile, do not scale to the quantities of text indexed by typical RAG systems. To combine the strengths of these contrasting methods, we propose GraphRAG, a graph-based approach to question answering over private text corpora that scales with both the generality of user questions and the quantity of source text. Our approach uses an LLM to build a graph index in two stages: first, to derive an entity knowledge graph from the source documents, then to pregenerate community summaries for all groups of closely related entities. Given a question, each community summary is used to generate a partial response, before all partial responses are again summarized in a final response to the user. For a class of global sensemaking questions over datasets in the 1 million token range, we show that', 'Large Language Model', 'Our approach uses an LLM to build a graph index in two stages: first, to derive an entity knowledge graph from the source documents, then to pregenerate community summaries for all groups of closely related entities.', "Despite their remarkable capabilities, large language models (LLMs) often produce responses containing factual inaccuracies due to their sole reliance on the parametric knowledge they encapsulate. Retrieval-Augmented Generation (RAG), an ad hoc approach that augments LMs with retrieval of relevant knowledge, decreases such issues. However, indiscriminately retrieving and incorporating a fixed number of retrieved passages, regardless of whether retrieval is necessary, or passages are relevant, diminishes LM versatility or can lead to unhelpful response generation. We introduce a new framework called Self-Reflective Retrieval-Augmented Generation (Self-RAG) that enhances an LM's quality and factuality through retrieval and self-reflection. Our framework trains a single arbitrary LM that adaptively retrieves passages on-demand, and generates and reflects on retrieved passages and its own generations using special tokens, called reflection tokens. Generating reflection tokens makes the LM controllable during the inference phase, enabling it to tailor its behavior to diverse task requirements. Experiments show that Self-RAG (7B and 13B parameters) significantly outperforms state-of-the-art LLMs and retrieval-augmented models on a diverse set of tasks. Specifically, Self-RAG outperforms ChatGPT and retrieval-augmented Llama2-chat on Open-domain QA, reasoning and fact verification tasks, and it shows significant gains in improving factuality and citation accuracy for long-form generations relative to these models.", 'Experiments show that Self-RAG (7B and 13B parameters) significantly outperforms state-of-the-art LLMs and retrieval-augmented models on a diverse set of tasks.', 'From RAG to Memory: Non-Parametric Continual Learning for Large Language Models', 'Our ability to continuously acquire, organize, and leverage knowledge is a key feature of human intelligence that AI systems must approximate to unlock their full potential. Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information. However, its reliance on vector retrieval hinders its ability to mimic the dynamic and interconnected nature of human long-term memory. Recent RAG approaches augment vector embeddings with various structures like knowledge graphs to address some of these gaps, namely sense-making and associativity. However, their performance on more basic factual memory tasks drops considerably below standard RAG. We address this unintended deterioration and propose HippoRAG 2, a framework that outperforms standard RAG comprehensively on factual, sense-making, and associative memory tasks. HippoRAG 2 builds upon the Personalized PageRank algorithm used in HippoRAG and enhances it with deeper passage integration and more effective online use of an LLM. This combination pushes this RAG system closer to the effectiveness of human long-term memory, achieving a 7% improvement in associative memory tasks over the state-of-the-art embedding model while also exhibiting superior factual knowledge and sense-making memory capabilities. This work paves the way for non-parametric continual learning', 'Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information.', 'Large Language Models (LLMs) have advanced artificial intelligence by enabling human-like text generation and natural language understanding. However, their reliance on static training data limits their ability to respond to dynamic, real-time queries, resulting in outdated or inaccurate outputs. Retrieval-Augmented Generation (RAG) has emerged as a solution, enhancing LLMs by integrating real-time data retrieval to provide contextually relevant and up-to-date responses. Despite its promise, traditional RAG systems are constrained by static workflows and lack the adaptability required for multi-step reasoning and complex task management. Agentic Retrieval-Augmented Generation (Agentic RAG) transcends these limitations by embedding autonomous AI agents into the RAG pipeline. These agents leverage agentic design patterns reflection, planning, tool use, and multi-agent collaboration to dynamically manage retrieval strategies, iteratively refine contextual understanding, and adapt workflows through operational structures ranging from sequential steps to adaptive collaboration. This integration enables Agentic RAG systems to deliver flexibility, scalability, and context-awareness across diverse applications. This paper presents an analytical survey of Agentic RAG systems. It traces the evolution of RAG paradigms, introduces a principled taxonomy of Agentic RAG architectures based on agent cardinality, control structure, autonomy, and knowledge representation, and provides a comparative analysis of design trade-offs across existing frameworks. The survey examines', 'Retrieval-Augmented Generation (RAG) overcomes the main barrier for the adoption of LLM-based chatbots in education: hallucinations. The uncomplicated architecture of RAG chatbots makes it relatively easy to implement chatbots that serve specific purposes and thus are capable of addressing various needs in the educational domain. With five years having passed since the introduction of RAG, the time has come to check the progress attained in its adoption in education. This paper identifies 47 papers dedicated to RAG chatbots’ uses for various kinds of educational purposes, which are analyzed in terms of their character, the target of the support provided by the chatbots, the thematic scope of the knowledge accessible via the chatbots, the underlying large language model, and the character of their evaluation.', 'Retrieval-Augmented Generation (RAG) overcomes the main barrier for the adoption of LLM-based chatbots in education: hallucinations.']

retrieval: ['The use of retrieval-augmented generation (RAG) to retrieve relevant information from an external knowledge source enables large language models (LLMs) to answer questions over private and/or previously unseen document collections. However, RAG fails on global questions directed at an entire text corpus, such as"What are the main themes in the dataset?", since this is inherently a query-focused summarization (QFS) task, rather than an explicit retrieval task. Prior QFS methods, meanwhile, do not scale to the quantities of text indexed by typical RAG systems. To combine the strengths of these contrasting methods, we propose GraphRAG, a graph-based approach to question answering over private text corpora that scales with both the generality of user questions and the quantity of source text. Our approach uses an LLM to build a graph index in two stages: first, to derive an entity knowledge graph from the source documents, then to pregenerate community summaries for all groups of closely related entities. Given a question, each community summary is used to generate a partial response, before all partial responses are again summarized in a final response to the user. For a class of global sensemaking questions over datasets in the 1 million token range, we show that', 'However, RAG fails on global questions directed at an entire text corpus, such as"What are the main themes in the dataset?", since this is inherently a query-focused summarization (QFS) task, rather than an explicit retrieval task.', 'Retrieval-Augmented Generation', "Despite their remarkable capabilities, large language models (LLMs) often produce responses containing factual inaccuracies due to their sole reliance on the parametric knowledge they encapsulate. Retrieval-Augmented Generation (RAG), an ad hoc approach that augments LMs with retrieval of relevant knowledge, decreases such issues. However, indiscriminately retrieving and incorporating a fixed number of retrieved passages, regardless of whether retrieval is necessary, or passages are relevant, diminishes LM versatility or can lead to unhelpful response generation. We introduce a new framework called Self-Reflective Retrieval-Augmented Generation (Self-RAG) that enhances an LM's quality and factuality through retrieval and self-reflection. Our framework trains a single arbitrary LM that adaptively retrieves passages on-demand, and generates and reflects on retrieved passages and its own generations using special tokens, called reflection tokens. Generating reflection tokens makes the LM controllable during the inference phase, enabling it to tailor its behavior to diverse task requirements. Experiments show that Self-RAG (7B and 13B parameters) significantly outperforms state-of-the-art LLMs and retrieval-augmented models on a diverse set of tasks. Specifically, Self-RAG outperforms ChatGPT and retrieval-augmented Llama2-chat on Open-domain QA, reasoning and fact verification tasks, and it shows significant gains in improving factuality and citation accuracy for long-form generations relative to these models.", 'Retrieval-Augmented Generation (RAG), an ad hoc approach that augments LMs with retrieval of relevant knowledge, decreases such issues.', "We introduce a new framework called Self-Reflective Retrieval-Augmented Generation (Self-RAG) that enhances an LM's quality and factuality through retrieval and self-reflection.", 'Our framework trains a single arbitrary LM that adaptively retrieves passages on-demand, and generates and reflects on retrieved passages and its own generations using special tokens, called reflection tokens.', 'Experiments show that Self-RAG (7B and 13B parameters) significantly outperforms state-of-the-art LLMs and retrieval-augmented models on a diverse set of tasks.', 'Our ability to continuously acquire, organize, and leverage knowledge is a key feature of human intelligence that AI systems must approximate to unlock their full potential. Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information. However, its reliance on vector retrieval hinders its ability to mimic the dynamic and interconnected nature of human long-term memory. Recent RAG approaches augment vector embeddings with various structures like knowledge graphs to address some of these gaps, namely sense-making and associativity. However, their performance on more basic factual memory tasks drops considerably below standard RAG. We address this unintended deterioration and propose HippoRAG 2, a framework that outperforms standard RAG comprehensively on factual, sense-making, and associative memory tasks. HippoRAG 2 builds upon the Personalized PageRank algorithm used in HippoRAG and enhances it with deeper passage integration and more effective online use of an LLM. This combination pushes this RAG system closer to the effectiveness of human long-term memory, achieving a 7% improvement in associative memory tasks over the state-of-the-art embedding model while also exhibiting superior factual knowledge and sense-making memory capabilities. This work paves the way for non-parametric continual learning', 'Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information.', 'Agentic Retrieval-Augmented Generation: A Survey on Agentic RAG', 'Large Language Models (LLMs) have advanced artificial intelligence by enabling human-like text generation and natural language understanding. However, their reliance on static training data limits their ability to respond to dynamic, real-time queries, resulting in outdated or inaccurate outputs. Retrieval-Augmented Generation (RAG) has emerged as a solution, enhancing LLMs by integrating real-time data retrieval to provide contextually relevant and up-to-date responses. Despite its promise, traditional RAG systems are constrained by static workflows and lack the adaptability required for multi-step reasoning and complex task management. Agentic Retrieval-Augmented Generation (Agentic RAG) transcends these limitations by embedding autonomous AI agents into the RAG pipeline. These agents leverage agentic design patterns reflection, planning, tool use, and multi-agent collaboration to dynamically manage retrieval strategies, iteratively refine contextual understanding, and adapt workflows through operational structures ranging from sequential steps to adaptive collaboration. This integration enables Agentic RAG systems to deliver flexibility, scalability, and context-awareness across diverse applications. This paper presents an analytical survey of Agentic RAG systems. It traces the evolution of RAG paradigms, introduces a principled taxonomy of Agentic RAG architectures based on agent cardinality, control structure, autonomy, and knowledge representation, and provides a comparative analysis of design trade-offs across existing frameworks. The survey examines', 'Retrieval-Augmented Generation (RAG) Chatbots for Education: A Survey of Applications', 'Retrieval-Augmented Generation (RAG) overcomes the main barrier for the adoption of LLM-based chatbots in education: hallucinations. The uncomplicated architecture of RAG chatbots makes it relatively easy to implement chatbots that serve specific purposes and thus are capable of addressing various needs in the educational domain. With five years having passed since the introduction of RAG, the time has come to check the progress attained in its adoption in education. This paper identifies 47 papers dedicated to RAG chatbots' uses for various kinds of educational purposes, which are analyzed in terms of their character, the target of the support provided by the chatbots, the thematic scope of the knowledge accessible via the chatbots, the underlying large language model, and the character of their evaluation.', 'Retrieval-Augmented Generation (RAG) overcomes the main barrier for the adoption of LLM-based chatbots in education: hallucinations.']

grounding: ["Despite their remarkable capabilities, large language models (LLMs) often produce responses containing factual inaccuracies due to their sole reliance on the parametric knowledge they encapsulate. Retrieval-Augmented Generation (RAG), an ad hoc approach that augments LMs with retrieval of relevant knowledge, decreases such issues. However, indiscriminately retrieving and incorporating a fixed number of retrieved passages, regardless of whether retrieval is necessary, or passages are relevant, diminishes LM versatility or can lead to unhelpful response generation. We introduce a new framework called Self-Reflective Retrieval-Augmented Generation (Self-RAG) that enhances an LM's quality and factuality through retrieval and self-reflection. Our framework trains a single arbitrary LM that adaptively retrieves passages on-demand, and generates and reflects on retrieved passages and its own generations using special tokens, called reflection tokens. Generating reflection tokens makes the LM controllable during the inference phase, enabling it to tailor its behavior to diverse task requirements. Experiments show that Self-RAG (7B and 13B parameters) significantly outperforms state-of-the-art LLMs and retrieval-augmented models on a diverse set of tasks. Specifically, Self-RAG outperforms ChatGPT and retrieval-augmented Llama2-chat on Open-domain QA, reasoning and fact verification tasks, and it shows significant gains in improving factuality and citation accuracy for long-form generations relative to these models.", "We introduce a new framework called Self-Reflective Retrieval-Augmented Generation (Self-RAG) that enhances an LM's quality and factuality through retrieval and self-reflection."]

hallucination: ['Retrieval-Augmented Generation (RAG) overcomes the main barrier for the adoption of LLM-based chatbots in education: hallucinations. The uncomplicated architecture of RAG chatbots makes it relatively easy to implement chatbots that serve specific purposes and thus are capable of addressing various needs in the educational domain. With five years having passed since the introduction of RAG, the time has come to check the progress attained in its adoption in education. This paper identifies 47 papers dedicated to RAG chatbots' uses for various kinds of educational purposes, which are analyzed in terms of their character, the target of the support provided by the chatbots, the thematic scope of the knowledge accessible via the chatbots, the underlying large language model, and the character of their evaluation.', 'Retrieval-Augmented Generation (RAG) overcomes the main barrier for the adoption of LLM-based chatbots in education: hallucinations.']

knowledge_freshness: ['From RAG to Memory: Non-Parametric Continual Learning for Large Language Models', 'Our ability to continuously acquire, organize, and leverage knowledge is a key feature of human intelligence that AI systems must approximate to unlock their full potential. Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information. However, its reliance on vector retrieval hinders its ability to mimic the dynamic and interconnected nature of human long-term memory. Recent RAG approaches augment vector embeddings with various structures like knowledge graphs to address some of these gaps, namely sense-making and associativity. However, their performance on more basic factual memory tasks drops considerably below standard RAG. We address this unintended deterioration and propose HippoRAG 2, a framework that outperforms standard RAG comprehensively on factual, sense-making, and associative memory tasks. HippoRAG 2 builds upon the Personalized PageRank algorithm used in HippoRAG and enhances it with deeper passage integration and more effective online use of an LLM. This combination pushes this RAG system closer to the effectiveness of human long-term memory, achieving a 7% improvement in associative memory tasks over the state-of-the-art embedding model while also exhibiting superior factual knowledge and sense-making memory capabilities. This work paves the way for non-parametric continual learning', 'Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information.']

evaluation: ["Despite their remarkable capabilities, large language models (LLMs) often produce responses containing factual inaccuracies due to their sole reliance on the parametric knowledge they encapsulate. Retrieval-Augmented Generation (RAG), an ad hoc approach that augments LMs with retrieval of relevant knowledge, decreases such issues. However, indiscriminately retrieving and incorporating a fixed number of retrieved passages, regardless of whether retrieval is necessary, or passages are relevant, diminishes LM versatility or can lead to unhelpful response generation. We introduce a new framework called Self-Reflective Retrieval-Augmented Generation (Self-RAG) that enhances an LM's quality and factuality through retrieval and self-reflection. Our framework trains a single arbitrary LM that adaptively retrieves passages on-demand, and generates and reflects on retrieved passages and its own generations using special tokens, called reflection tokens. Generating reflection tokens makes the LM controllable during the inference phase, enabling it to tailor its behavior to diverse task requirements. Experiments show that Self-RAG (7B and 13B parameters) significantly outperforms state-of-the-art LLMs and retrieval-augmented models on a diverse set of tasks. Specifically, Self-RAG outperforms ChatGPT and retrieval-augmented Llama2-chat on Open-domain QA, reasoning and fact verification tasks, and it shows significant gains in improving factuality and citation accuracy for long-form generations relative to these models.", 'Experimental Evaluation', 'From RAG to Memory: Non-Parametric Continual Learning for Large Language Models', 'Our ability to continuously acquire, organize, and leverage knowledge is a key feature of human intelligence that AI systems must approximate to unlock their full potential. Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information. However, its reliance on vector retrieval hinders its ability to mimic the dynamic and interconnected nature of human long-term memory. Recent RAG approaches augment vector embeddings with various structures like knowledge graphs to address some of these gaps, namely sense-making and associativity. However, their performance on more basic factual memory tasks drops considerably below standard RAG. We address this unintended deterioration and propose HippoRAG 2, a framework that outperforms standard RAG comprehensively on factual, sense-making, and associative memory tasks. HippoRAG 2 builds upon the Personalized PageRank algorithm used in HippoRAG and enhances it with deeper passage integration and more effective online use of an LLM. This combination pushes this RAG system closer to the effectiveness of human long-term memory, achieving a 7% improvement in associative memory tasks over the state-of-the-art embedding model while also exhibiting superior factual knowledge and sense-making memory capabilities. This work paves the way for non-parametric continual learning', 'Retrieval-Augmented Generation (RAG) overcomes the main barrier for the adoption of

LLM-based chatbots in education: hallucinations. The uncomplicated architecture of RAG chatbots makes it relatively easy to implement chatbots that serve specific purposes and thus are capable of addressing various needs in the educational domain. With five years having passed since the introduction of RAG, the time has come to check the progress attained in its adoption in education. This paper identifies 47 papers dedicated to RAG chatbots' uses for various kinds of educational purposes, which are analyzed in terms of their character, the target of the support provided by the chatbots, the thematic scope of the knowledge accessible via the chatbots, the underlying large language model, and the character of their evaluation.']

applications: ['Large Language Models (LLMs) have advanced artificial intelligence by enabling human-like text generation and natural language understanding. However, their reliance on static training data limits their ability to respond to dynamic, real-time queries, resulting in outdated or inaccurate outputs. Retrieval-Augmented Generation (RAG) has emerged as a solution, enhancing LLMs by integrating real-time data retrieval to provide contextually relevant and up-to-date responses. Despite its promise, traditional RAG systems are constrained by static workflows and lack the adaptability required for multi-step reasoning and complex task management. Agentic Retrieval-Augmented Generation (Agentic RAG) transcends these limitations by embedding autonomous AI agents into the RAG pipeline. These agents leverage agentic design patterns reflection, planning, tool use, and multi-agent collaboration to dynamically manage retrieval strategies, iteratively refine contextual understanding, and adapt workflows through operational structures ranging from sequential steps to adaptive collaboration. This integration enables Agentic RAG systems to deliver flexibility, scalability, and context-awareness across diverse applications. This paper presents an analytical survey of Agentic RAG systems. It traces the evolution of RAG paradigms, introduces a principled taxonomy of Agentic RAG architectures based on agent cardinality, control structure, autonomy, and knowledge representation, and provides a comparative analysis of design trade-offs across existing frameworks. The survey examines', 'Retrieval-Augmented Generation (RAG) Chatbots for Education: A Survey of Applications']

coordination: ['Large Language Models (LLMs) have advanced artificial intelligence by enabling human-like text generation and natural language understanding. However, their reliance on static training data limits their ability to respond to dynamic, real-time queries, resulting in outdated or inaccurate outputs. Retrieval-Augmented Generation (RAG) has emerged as a solution, enhancing LLMs by integrating real-time data retrieval to provide contextually relevant and up-to-date responses. Despite its promise, traditional RAG systems are constrained by static workflows and lack the adaptability required for multi-step reasoning and complex task management. Agentic Retrieval-Augmented Generation (Agentic RAG) transcends these limitations by embedding autonomous AI agents into the RAG pipeline. These agents leverage agentic design patterns reflection, planning, tool use, and multi-agent collaboration to dynamically manage retrieval strategies, iteratively refine contextual understanding, and adapt workflows through operational structures ranging from sequential steps to adaptive collaboration. This integration enables Agentic RAG systems to deliver flexibility, scalability, and context-awareness across diverse applications. This paper presents an analytical survey of Agentic RAG systems. It traces the evolution of RAG paradigms, introduces a principled taxonomy of Agentic RAG architectures based on agent cardinality, control structure, autonomy, and knowledge representation, and provides a comparative analysis of design trade-offs across existing frameworks. The survey examines']

security: []

scalability: ['The use of retrieval-augmented generation (RAG) to retrieve relevant information from an external knowledge source enables large language models (LLMs) to answer questions over private and/or previously unseen document collections. However, RAG fails on global questions directed at an entire text corpus, such as"What are the main themes in the dataset?", since this is inherently a query-focused summarization (QFS) task, rather than an explicit retrieval task. Prior QFS methods, meanwhile, do not scale to the quantities of text indexed by typical RAG systems. To combine the strengths of these contrasting methods, we propose GraphRAG, a graph-based approach to question answering over private text corpora that scales with both the generality of user questions and the quantity of source text. Our approach uses an LLM to build a graph index in two stages: first, to derive an entity knowledge graph from the source documents, then to pregenerate community summaries for all groups of closely related entities. Given a question, each community summary is used to generate a partial response, before all partial responses are again summarized in a final response to the user. For a class of global sensemaking questions over datasets in the 1 million token range, we show that', 'To combine the strengths of these contrasting methods, we propose GraphRAG, a graph-based approach to question answering over private text corpora that scales with both the generality of user questions and the quantity of source text.', 'Large Language Models (LLMs) have advanced artificial intelligence by enabling human-like text generation and natural language understanding. However, their reliance on static training data limits their ability to respond to dynamic, real-time queries, resulting in outdated or inaccurate outputs. Retrieval-Augmented Generation (RAG) has emerged as a solution, enhancing LLMs by integrating real-time data retrieval to provide contextually relevant and up-to-date responses. Despite its promise, traditional RAG systems are constrained by static workflows and lack the adaptability required for multi-step reasoning and complex task management. Agentic Retrieval-Augmented Generation (Agentic RAG) transcends these limitations by embedding autonomous AI agents into the RAG pipeline. These agents leverage agentic design patterns reflection, planning, tool use, and multi-agent collaboration to dynamically manage retrieval strategies, iteratively refine contextual understanding, and adapt workflows through operational structures ranging from sequential steps to adaptive collaboration. This integration enables Agentic RAG systems to deliver flexibility, scalability, and context-awareness across diverse applications. This paper presents an analytical survey of Agentic RAG systems. It traces the evolution of RAG paradigms, introduces a principled taxonomy of Agentic RAG architectures based on agent cardinality, control structure, autonomy, and knowledge representation, and provides a comparative analysis of design trade-offs across existing frameworks. The survey examines']

explainability: []

reasoning: ["Despite their remarkable capabilities, large language models (LLMs) often produce responses containing factual inaccuracies due to their sole reliance on the parametric knowledge they encapsulate. Retrieval-Augmented Generation (RAG), an ad hoc approach that augments LMs with retrieval of relevant knowledge, decreases such issues. However, indiscriminately retrieving and incorporating a fixed number of retrieved passages, regardless of whether retrieval is necessary, or passages are relevant, diminishes LM versatility or can lead to unhelpful response generation. We introduce a new framework called Self-Reflective Retrieval-Augmented Generation (Self-RAG) that enhances an LM's quality and factuality through retrieval and self-reflection. Our framework trains a single arbitrary LM that adaptively retrieves passages on-demand, and generates and reflects on retrieved passages and its own generations using special tokens, called reflection tokens. Generating reflection tokens makes the LM controllable during the inference phase, enabling it to tailor its behavior to diverse task requirements. Experiments show that Self-RAG (7B and 13B parameters) significantly outperforms state-of-the-art LLMs and retrieval-augmented models on a diverse set of tasks. Specifically, Self-RAG outperforms ChatGPT and retrieval-augmented Llama2-chat on Open-domain QA, reasoning and fact verification tasks, and it shows significant gains in improving factuality and citation accuracy for long-form generations relative to these models.", 'Large Language Models (LLMs) have advanced artificial intelligence by enabling human-like text generation and natural language understanding. However, their reliance on static training data limits their ability to respond to dynamic, real-time queries, resulting in outdated or inaccurate outputs. Retrieval-Augmented Generation (RAG) has emerged as a solution, enhancing LLMs by integrating real-time data retrieval to provide contextually relevant and up-to-date responses. Despite its promise, traditional RAG systems are constrained by static workflows and lack the adaptability required for multi-step reasoning and complex task management. Agentic Retrieval-Augmented Generation (Agentic RAG) transcends these limitations by embedding autonomous AI agents into the RAG pipeline. These agents leverage agentic design patterns reflection, planning, tool use, and multi-agent collaboration to dynamically manage retrieval strategies, iteratively refine contextual understanding, and adapt workflows through operational structures ranging from sequential steps to adaptive collaboration. This integration enables Agentic RAG systems to deliver flexibility, scalability, and context-awareness across diverse applications. This paper presents an analytical survey of Agentic RAG systems. It traces the evolution of RAG paradigms, introduces a principled taxonomy of Agentic RAG architectures based on agent cardinality, control structure, autonomy, and knowledge representation, and provides a comparative analysis of design trade-offs across existing frameworks. The survey examines'] memory: ['Our ability to continuously acquire, organize, and leverage knowledge is a key feature of human intelligence that AI systems must approximate to unlock their full potential. Given the challenges in continual learning with large language models (LLMs), retrieval-augmented generation (RAG) has become the dominant way to introduce new information. However, its reliance on vector retrieval hinders its ability to mimic the dynamic and interconnected nature of human long-term memory. Recent RAG approaches augment vector embeddings with various structures like knowledge graphs to address some of these gaps, namely sense-making and associativity. However, their performance on more basic factual memory tasks drops considerably below standard RAG. We address this unintended deterioration and propose HippoRAG 2, a framework that outperforms standard RAG comprehensively on factual, sense-making, and associative memory tasks. HippoRAG 2 builds upon the Personalized PageRank algorithm used in HippoRAG and enhances it with deeper passage integration and more effective online use of an LLM. This combination pushes this RAG system closer to the effectiveness of human long-term memory, achieving a 7% improvement in associative memory tasks over the state-of-the-art embedding model while also exhibiting superior factual knowledge and sense-making memory capabilities. This work paves the way for non-parametric continual learning']

planning: ['Large Language Models (LLMs) have advanced artificial intelligence by enabling human-like text generation and natural language understanding. However, their reliance on static training data limits their ability to respond to dynamic, real-time queries, resulting in outdated or inaccurate outputs. Retrieval-Augmented Generation (RAG) has emerged as a solution, enhancing LLMs by integrating real-time data retrieval to provide contextually relevant and up-to-date responses. Despite its promise, traditional RAG systems are constrained by static workflows and lack the adaptability required for multi-step reasoning and complex task management. Agentic Retrieval-Augmented Generation (Agentic RAG) transcends these limitations by embedding autonomous AI agents into the RAG pipeline. These agents leverage agentic design patterns reflection, planning, tool use, and multi-agent collaboration to dynamically manage retrieval strategies, iteratively refine contextual understanding, and adapt workflows through operational structures ranging from sequential steps to adaptive collaboration. This integration enables Agentic RAG systems to deliver flexibility, scalability, and context-awareness across diverse applications. This paper presents an analytical survey of Agentic RAG systems. It traces the evolution of RAG paradigms, introduces a principled taxonomy of Agentic RAG architectures based on agent cardinality, control structure, autonomy, and knowledge representation, and provides a comparative analysis of design trade-offs across existing frameworks. The survey examines']

general: []

Research Gaps

• Retrieval quality and the ability to consistently retrieve relevant information for complex queries remain important research challenges.

• Improved grounding, factuality, and reliable source attribution are required to reduce unsupported or incorrect generated information.

• Hallucination and factual reliability remain unresolved challenges in current AI systems.

• Maintaining fresh and continuously updated knowledge remains a challenge for systems operating over changing information sources.

• More robust and standardized evaluation and benchmarking methods are required.

• Larger, more diverse, and domain-specific datasets are needed for reliable development and evaluation.

• Current models and architectures still require improvements in reliability, adaptability, and generalization.

• Reliable coordination and communication between multiple agents remains an open research challenge.

• Scalability, computational cost, and efficient resource usage remain challenges for practical deployment.

• Current systems require stronger reasoning capabilities for complex multi-step tasks.

Emerging Topics

• knowledge

• retrieval-augmented

• generation

• graph

• reasoning

• state-of-the-art

Recommendations

• Explore more reliable autonomous and multi-agent systems with improved coordination and task execution.

• Investigate improved inter-agent communication and coordination mechanisms for reliable multi-agent task execution.

• Improve reasoning capabilities through structured planning, multi-step inference, and verification mechanisms.

• Investigate more reliable autonomous planning and task decomposition techniques.

• Investigate long-term memory mechanisms for maintaining context across complex research and reasoning tasks.

• Improve retrieval quality and document selection for complex research queries.

• Strengthen source grounding and citation verification to improve factual reliability.

• Develop stronger hallucination detection and factual verification mechanisms.

• Investigate mechanisms for continuously updating knowledge from recent information sources.

• Evaluate proposed approaches using standardized benchmarks and reproducible evaluation protocols.

Recommendation Validation

• {'recommendation': 'Explore more reliable autonomous and multi-agent systems with improved coordination and task execution.', 'validation_status': 'Partially Supported', 'evidence_score': 1, 'evidence': {'research_gaps': True, 'future_work': False, 'emerging_topics': False, 'methodology': False}}

• {'recommendation': 'Investigate improved inter-agent communication and coordination mechanisms for reliable multi-agent task execution.', 'validation_status': 'Partially Supported', 'evidence_score': 1, 'evidence': {'research_gaps': True, 'future_work': False, 'emerging_topics': False, 'methodology': False}}

• {'recommendation': 'Improve reasoning capabilities through structured planning, multi-step inference, and verification mechanisms.', 'validation_status': 'Supported', 'evidence_score': 2, 'evidence': {'research_gaps': True, 'future_work': False, 'emerging_topics': True, 'methodology': False}}

• {'recommendation': 'Investigate more reliable autonomous planning and task decomposition techniques.', 'validation_status': 'Weakly Supported', 'evidence_score': 0, 'evidence': {'research_gaps': False, 'future_work': False, 'emerging_topics': False, 'methodology': False}}

- {'recommendation': 'Investigate long-term memory mechanisms for maintaining context across complex research and reasoning tasks.', 'validation_status': 'Strongly Supported', 'evidence_score': 3, 'evidence': {'research_gaps': True, 'future_work': False, 'emerging_topics': True, 'methodology': True}}
- {'recommendation': 'Improve retrieval quality and document selection for complex research queries.', 'validation_status': 'Strongly Supported', 'evidence_score': 3, 'evidence': {'research_gaps': True, 'future_work': False, 'emerging_topics': True, 'methodology': True}}
- {'recommendation': 'Strengthen source grounding and citation verification to improve factual reliability.', 'validation_status': 'Partially Supported', 'evidence_score': 1, 'evidence': {'research_gaps': True, 'future_work': False, 'emerging_topics': False, 'methodology': False}}
- {'recommendation': 'Develop stronger hallucination detection and factual verification mechanisms.', 'validation_status': 'Supported', 'evidence_score': 2, 'evidence': {'research_gaps': True, 'future_work': False, 'emerging_topics': False, 'methodology': True}}
- {'recommendation': 'Investigate mechanisms for continuously updating knowledge from recent information sources.', 'validation_status': 'Supported', 'evidence_score': 2, 'evidence': {'research_gaps': True, 'future_work': False, 'emerging_topics': True, 'methodology': False}}
- {'recommendation': 'Evaluate proposed approaches using standardized benchmarks and reproducible evaluation protocols.', 'validation_status': 'Supported', 'evidence_score': 2, 'evidence': {'research_gaps': True, 'future_work': False, 'emerging_topics': False, 'methodology': True}}

Summary

dominant_area: Retrieval-Augmented Generation

top_trend: knowledge

future_work_items: 0

research_gap_count: 10

validated_recommendation_count: 10

Experiment Plan

Research Areas

- Retrieval-Augmented Generation

Recommended Datasets

- BEIR
- HotpotQA
- MS MARCO
- Natural Questions
- TriviaQA

Baseline Models

- BM25 Retrieval
- Dense Retrieval RAG
- Hybrid Search RAG

- Naive RAG
- Vanilla RAG

Evaluation Metrics

- Accuracy
- Precision
- Recall
- F1-Score
- Retrieval Precision
- Retrieval Recall
- Context Relevance
- Answer Relevance
- Faithfulness
- Groundedness
- ROC-AUC

Hardware Requirements

cpu: Intel Core i7 / AMD Ryzen 7

ram: 32 GB

gpu: NVIDIA RTX 3060 or better

storage: 200 GB SSD

Validation Strategy

- Train-Test Split
- K-Fold Cross Validation
- Hyperparameter Tuning
- Statistical Significance Testing
- Retrieval Evaluation
- Answer Quality Evaluation

Experimental Workflow

- Collect Knowledge Documents
- Preprocess and Chunk Documents
- Build Document Index
- Configure Retrieval System
- Retrieve Relevant Context
- Generate Grounded Response
- Evaluate Retrieval Quality
- Evaluate Answer Quality

- Compare Baseline and Proposed RAG
- Analyze Errors
- Draw Conclusions

Gap Based Recommendations

- Evaluate the proposed approach on larger and more diverse datasets.
- Compare BM25, dense retrieval, and hybrid retrieval methods.
- Evaluate source attribution, faithfulness, and groundedness.
- Measure hallucination frequency and introduce factuality verification.
- Evaluate system performance using recently updated knowledge sources.
- Use standardized benchmarks and multiple evaluation metrics.
- Compare centralized and decentralized multi-agent coordination strategies.
- Measure computational cost, latency, resource consumption, and scalability.
- Evaluate multi-step reasoning using complex benchmark tasks.
- Evaluate long-term memory and persistent context management.
- Evaluate long-horizon planning and decision-making performance.
- Validate the system in realistic real-world deployment scenarios.

Report Summary

Dominant Research Area: Retrieval-Augmented Generation

Top Research Trend: knowledge

Highest Cited Paper: Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection

Latest Paper: From RAG to Memory: Non-Parametric Continual Learning for Large Language Models

Recommended Datasets

- BEIR
- HotpotQA
- MS MARCO
- Natural Questions
- TriviaQA

Baseline Models

- BM25 Retrieval
- Dense Retrieval RAG
- Hybrid Search RAG
- Naive RAG
- Vanilla RAG

Evaluation Metrics

- Accuracy
- Precision
- Recall
- F1-Score
- Retrieval Precision
- Retrieval Recall
- Context Relevance
- Answer Relevance
- Faithfulness
- Groundedness
- ROC-AUC

Research Recommendations

- Explore more reliable autonomous and multi-agent systems with improved coordination and task execution.
- Investigate improved inter-agent communication and coordination mechanisms for reliable multi-agent task execution.
- Improve reasoning capabilities through structured planning, multi-step inference, and verification mechanisms.
- Investigate more reliable autonomous planning and task decomposition techniques.
- Investigate long-term memory mechanisms for maintaining context across complex research and reasoning tasks.
- Improve retrieval quality and document selection for complex research queries.
- Strengthen source grounding and citation verification to improve factual reliability.
- Develop stronger hallucination detection and factual verification mechanisms.
- Investigate mechanisms for continuously updating knowledge from recent information sources.
- Evaluate proposed approaches using standardized benchmarks and reproducible evaluation protocols.

Research Gap Count: 10

Future Work Items: 0

Conclusion

The report has been generated using a personalized multi-agent pipeline. The explanation style was adapted according to the Intermediate profile to improve readability and usefulness.